

RAG Chatbot Sample Document

Introduction to RAG

Retrieval-Augmented Generation (RAG) is a technique used to give LLMs access to specific, external data without needing to retrain the model. It combines the power of retrieval systems with generative AI.

How it Works

1. Data Ingestion: Documents are loaded, split into chunks, and embedded into a vector store.
2. Retrieval: When a user asks a question, the system finds the most relevant chunks.
3. Generation: The LLM uses the retrieved chunks as context to generate an accurate answer.

Benefits

- Accuracy: Reduces hallucinations by providing factual context.
- Up-to-date: Can access real-time data or internal company documents.
- Efficiency: Faster and cheaper than fine-tuning models.